

Hidden Markov Models

Theoretical Aspects. Octave Implementation. Applications

Alexandru Sorici, Tudor Berariu

Romanian Asociation for Artificial Intelligence

October, 27th, 2012

PART 1.

Intro

PART 2.

Theory of HMMs

PART 3.

Demo & Discussions

Outline

Intro

1 ARIA Education Workshops

- ARIA's Mission
- ARIA Education
- Workshop Program

1 ARIA Education Workshops

- ARIA's Mission
- ARIA Education
- Workshop Program

1 ARIA Education Workshops

- ARIA's Mission
- ARIA Education
- Workshop Program

ARIA EDU

:)

1 ARIA Education Workshops

- ARIA's Mission
- ARIA Education
- Workshop Program

Today's Program

9:00	Registration
10:00	Rahaturi despre ARIA
11:00	HMM Theory

Theory of HMMs

2 Machine Learning Applications for HMM

- Machine Learning
- Where do HMMs fit into Machine Learning?

3 Theory of HMMs

- The 3 things you want from an HMM
- Mathematical Foundations for HMMs
- Notation Conventions & Framework Description

4 Implementing HMMs

- Using the Model for Estimations: the Forward-Backward algorithm
- Learning from Observations: Baum-Welch algorithm
- Uncovering Hidden states: Viterbi algorithm

2 Machine Learning Applications for HMM

- Machine Learning
- Where do HMMs fit into Machine Learning?

3 Theory of HMMs

- The 3 things you want from an HMM
- Mathematical Foundations for HMMs
- Notation Conventions & Framework Description

4 Implementing HMMs

- Using the Model for Estimations: the Forward-Backward algorithm
- Learning from Observations: Baum-Welch algorithm
- Uncovering Hidden states: Viterbi algorithm

What is Machine Learning?

Machine Learning

..**Trascau** is a beautiful horse.

Machine Learning Applications

- Computer Vision: Google Car
- Machine Translation
- Speech Recognition
- Recommender Systems
- Intelligent Advertising

2 Machine Learning Applications for HMM

- Machine Learning
- Where do HMMs fit into Machine Learning?

3 Theory of HMMs

- The 3 things you want from an HMM
- Mathematical Foundations for HMMs
- Notation Conventions & Framework Description

4 Implementing HMMs

- Using the Model for Estimations: the Forward-Backward algorithm
- Learning from Observations: Baum-Welch algorithm
- Uncovering Hidden states: Viterbi algorithm

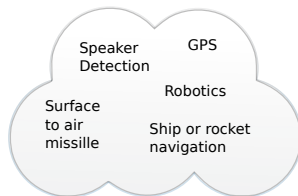
Machine Learning Classification

Types of Machine Learning Problems


- Regression
- Classification
- Reinforcement Learning
- supervised learning (eg. ..)
- unsupervised

Sequence / Temporal problems (I)

OBJECT TRACKING



SPEECH RECOGNITION



GESTURE RECOGNITION



Sequence / Temporal problems (II)

BIOINFORMATICS

Protein Sequencing

Modeling of a Gene
Regulatory Network

ECONOMICS

Stock Price Prediction

Econometrics

- estimate a country's economic indicators across time -

Probabilistic Reasoning over Time - Models

Consider some of the previously presented problems ...

Probabilistic Reasoning over Time - Models

Consider some of the previously presented problems ...

How do we model such dynamic situations?

Probabilistic Reasoning over Time - Models

Consider some of the previously presented problems ...

How do we model such dynamic situations?

States and Observations

- The process of change is viewed as a series of **time slices (snapshots)**
- Each time slice contains a set of random variables
 - \mathbf{O}_t - set of all **observable** evidence variables at time t
 - \mathbf{Q}_t - set of all **unobservable / hidden** state variables at time t

Probabilistic Reasoning over Time - Assumptions

Consider some of the previously presented problems ...

Probabilistic Reasoning over Time - Assumptions

Consider some of the previously presented problems ...

What **assumptions** (if any) do we make?

Probabilistic Reasoning over Time - Assumptions

Consider some of the previously presented problems ...

What **assumptions** (if any) do we make?

Stationary Process

The process of change is governed by laws **that do not themselves change over time**.

Implication: we need to specify conditional distributions only for the variables within a *representative* timeslice.

Probabilistic Reasoning over Time - Assumptions

Consider some of the previously presented problems ...

What **assumptions** (if any) do we make?

Stationary Process

The process of change is governed by laws **that do not themselves change over time**.

Implication: we need to specify conditional distributions only for the variables within a *representative* timeslice.

Markov Assumption

The current state in a process of change depends only on a **finite history** of previous states.

Implication: there is a **bounded** number of “parents” for the variables in each time slice.

$$P(Q_t | Q_{1:t-1}) = P(Q_t | Q_{t-1}) \quad P(O_t | Q_{1:t}, Q_{1:t-1}) = P(O_t | Q_t)$$

Probabilistic Reasoning over Time - Inference

What are the basic inference tasks that must be solved?

Probabilistic Reasoning over Time - Inference

What are the basic inference tasks that must be solved?

Filtering (monitoring)

The task of computing the **belief state** - the posterior distribution over the **current state**, given all evidence to date.

$$P(\mathbf{Q}_t | \mathbf{o}_{1:t})$$

Probabilistic Reasoning over Time - Inference

What are the basic inference tasks that must be solved?

Filtering (monitoring)

The task of computing the **belief state** - the posterior distribution over the **current state**, given all evidence to date.

$$P(\mathbf{Q}_t | \mathbf{o}_{1:t})$$

Evaluation (likelihood)

The task of computing the **likelihood** of the evidence up to present.

$$P(\mathbf{o}_{1:t})$$

Probabilistic Reasoning over Time - Inference

Prediction

The task of computing the posterior distribution over the **future state**, given all evidence to date.

$P(\mathbf{Q}_{t+k} | \mathbf{o}_{1:t})$, for some $k > 0$

Probabilistic Reasoning over Time - Inference

Prediction

The task of computing the posterior distribution over the **future state**, given all evidence to date.

$$P(\mathbf{Q}_{t+k} | \mathbf{o}_{1:t}), \text{ for some } k > 0$$

Smoothing (hindsight)

The task of computing the posterior distribution over a **past state**, given all evidence to the present.

$$P(\mathbf{Q}_k | \mathbf{o}_{1:t}), \text{ for some } 1 \leq k < t$$

Provides a better estimate of the state than was available at the time.

Probabilistic Reasoning over Time - Inference

Most likely explanation

Given a *sequence of observations*, find the **sequence of states** that is **most likely** to have generated those observations. $\operatorname{argmax}_{q_{1:t}} \mathbf{P}(\mathbf{q}_{t+k} | \mathbf{o}_{1:t})$, for some $k > 0$

Probabilistic Reasoning over Time - Inference

Most likely explanation

Given a *sequence of observations*, find the **sequence of states** that is **most likely** to have generated those observations. $\operatorname{argmax}_{q_{1:t}} \mathbf{P}(\mathbf{q}_{t+k} | \mathbf{o}_{1:t})$, for some $k > 0$

Learning

Given a set of *observation sequences*, find a method to learn the **transition** (e.g. $\mathbf{P}(\mathbf{q}_{t+1} = s_j | \mathbf{q}_t = s_i)$, $1 \leq i, j < N$) and **sensor** ($\mathbf{P}(\mathbf{o}_t | \mathbf{q}_t)$) **models** from the observations.

Probabilistic Reasoning over Time - Known Methods

Dynamic Bayesian Networks (DBN)

A DBN is Bayesian network that represents a temporal probability model.

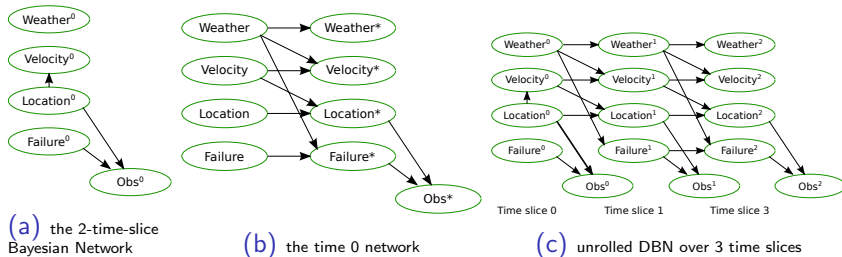



Figure: A highly simplified DBN for monitoring a vehicle (Koller and Friedman 2009)

Applied in problems like: object tracking, human activity recognition, protein sequencing etc.

Probabilistic Reasoning over Time - Known Methods

Kalman Filters (Linear Dynamical Systems)

A temporal model of one or more real-valued variables that **evolve linearly** over time, with some **Gaussian noise**.



- can be viewed as DBNs where all variables are continuous and all dependencies are linear gaussian
- wide application in **object tracking**

Figure: BN structure for a linear dynamical system with position X_t , velocity \dot{X}_t , and position measurement Z_t

Probabilistic Reasoning over Time - Known Methods

Hidden Markov Models (HMM)

An HMM is a temporal probabilistic model in which the state of the process is described by a **single discrete** random variable. The possible values of the variable are the possible states of the world.

Used successfully in applications like:

- Handwriting Recognition
- Gesture Recognition
- Speech Recognition
- Part-of-Speech Tagging
- DNA Sequencing

2 Machine Learning Applications for HMM

- Machine Learning
- Where do HMMs fit into Machine Learning?

3 Theory of HMMs

- The 3 things you want from an HMM
- Mathematical Foundations for HMMs
- Notation Conventions & Framework Description

4 Implementing HMMs

- Using the Model for Estimations: the Forward-Backward algorithm
- Learning from Observations: Baum-Welch algorithm
- Uncovering Hidden states: Viterbi algorithm

The 3 fundamental problems (Rabiner 1989)

- Particularization of temporal inference problems to the HMM case
- The restricted structure of the HMM allows for elegant implementations of all the basic algorithms

The 3 fundamental problems (Rabiner 1989)

- Particularization of temporal inference problems to the HMM case
- The restricted structure of the HMM allows for elegant implementations of all the basic algorithms

Evaluation Problem

Given a model and a sequence of observations, how do we compute the probability that the **observed sequence** was produced by the model?

The 3 fundamental problems (Rabiner 1989)

- Particularization of temporal inference problems to the HMM case
- The restricted structure of the HMM allows for elegant implementations of all the basic algorithms

Evaluation Problem

Given a model and a sequence of observations, how do we compute the probability that the **observed sequence** was produced by the model?

Best Explanation of Observations Problem

Given a model and a sequence of observations how do we choose a corresponding sequence of **states** which *gives meaning* to the observations?
How do we *uncover* the hidden part of the model?

The 3 fundamental problems (Rabiner 1989)

- Particularization of temporal inference problems to the HMM case
- The restricted structure of the HMM allows for elegant implementations of all the basic algorithms

Evaluation Problem

Given a model and a sequence of observations, how do we compute the probability that the **observed sequence** was produced by the model?

Best Explanation of Observations Problem

Given a model and a sequence of observations how do we choose a corresponding sequence of **states** which *gives meaning* to the observations?
How do we *uncover* the hidden part of the model?

Model Estimation (Training) Problem

Given some observed sequences, how do we adjust the **parameters** of an HMM model that best tries to explain the observations?

Outline

2 Machine Learning Applications for HMM

- Machine Learning
- Where do HMMs fit into Machine Learning?

3 Theory of HMMs

- The 3 things you want from an HMM
- **Mathematical Foundations for HMMs**
- Notation Conventions & Framework Description

4 Implementing HMMs

- Using the Model for Estimations: the Forward-Backward algorithm
- Learning from Observations: Baum-Welch algorithm
- Uncovering Hidden states: Viterbi algorithm

Elements of an HMM

- N Hidden States : $S_1, S_2, \dots S_N$
- M Observable Variables : $O_1, O_2, \dots O_M$

Parameters:

- Transition Function / Matrix between states
- Emission probabilities
- Initial state probabilities

Formalisation of the estimation problem



Formalisation of problem # 2

- $P(Q_1) = \sum_{x=\{1,2\}}^N P(Q_2)\theta\Pi$
- $P(Q_i|q_i = s_x) = i \times x \cdot i \dots$

Formalisation of parameters estimation problem



2 Machine Learning Applications for HMM

- Machine Learning
- Where do HMMs fit into Machine Learning?

3 Theory of HMMs

- The 3 things you want from an HMM
- Mathematical Foundations for HMMs
- **Notation Conventions & Framework Description**

4 Implementing HMMs

- Using the Model for Estimations: the Forward-Backward algorithm
- Learning from Observations: Baum-Welch algorithm
- Uncovering Hidden states: Viterbi algorithm

Notation Conventions

Variables in Octave

Outline

2 Machine Learning Applications for HMM

- Machine Learning
- Where do HMMs fit into Machine Learning?

3 Theory of HMMs

- The 3 things you want from an HMM
- Mathematical Foundations for HMMs
- Notation Conventions & Framework Description

4 Implementing HMMs

- Using the Model for Estimations: the Forward-Backward algorithm
- Learning from Observations: Baum-Welch algorithm
- Uncovering Hidden states: Viterbi algorithm

Outline

2 Machine Learning Applications for HMM

- Machine Learning
- Where do HMMs fit into Machine Learning?

3 Theory of HMMs

- The 3 things you want from an HMM
- Mathematical Foundations for HMMs
- Notation Conventions & Framework Description

4 Implementing HMMs

- Using the Model for Estimations: the Forward-Backward algorithm
- Learning from Observations: Baum-Welch algorithm
- Uncovering Hidden states: Viterbi algorithm

Learning from observations - Reminder

Model Estimation (Training) Problem

Given some observed sequences, how do we adjust the **parameters** of an HMM model that best tries to explain the observations?

Learning from observations - Reminder

Model Estimation (Training) Problem

Given some observed sequences, how do we adjust the **parameters** of an HMM model that best tries to explain the observations?

Adjust the model parameters $\lambda = (A, B, \Pi)$ to obtain $\max_{\lambda} P(O|\lambda)$

Learning from observations - Reminder

Model Estimation (Training) Problem

Given some observed sequences, how do we adjust the **parameters** of an HMM model that best tries to explain the observations?

Adjust the model parameters $\lambda = (A, B, \Pi)$ to obtain $\max_{\lambda} P(O|\lambda)$

The observation sequence used to adjust the model parameters is called a **training** sequence.

Training problem is crucial - allows to create best models for real phenomena.

Learning from observations - Aspects of the approach

Learning from observations - Aspects of the approach

Problem

There is no known way to analytically solve for the model which maximizes the probability of the observation sequence.

Learning from observations - Aspects of the approach

Problem

There is no known way to analytically solve for the model which maximizes the probability of the observation sequence.

Solution

We can choose $\lambda = (A, B, \Pi)$ such that $\max_{\lambda} P(O|\lambda)$ is **locally maximized** using an **iterative procedure** such as *Baum-Welch*.

The method is an instance of the *EM algorithm* (Dempster, Laird, and Rubin 1977) for HMMs.

Baum-Welch algorithm (I)

Baum-Welch algorithm (I)

We first define some auxiliary values:

$$\xi_{t,i,j} = \xi_t(i,j) = P(q_t = s_i, q_{t+1} = s_j | O, \lambda)$$

The probability of being in state s_i at time t and in state s_j at time $t + 1$, given the model and the observation sequence.

Baum-Welch algorithm (I)

We first define some auxiliary values:

$$\xi_{t,i,j} = \xi_t(i,j) = P(q_t = s_i, q_{t+1} = s_j | O, \lambda)$$

The probability of being in state s_i at time t and in state s_j at time $t + 1$, given the model and the observation sequence.

$$\gamma_{t,i} = \gamma_t(i) = P(q_t = s_i | O, \lambda)$$

The probability of being in state s_i at time t , given the model and the observation sequence.

Baum-Welch algorithm (I)

We first define some auxiliary values:

$$\xi_{t,i,j} = \xi_t(i,j) = P(q_t = s_i, q_{t+1} = s_j | O, \lambda)$$

The probability of being in state s_i at time t and in state s_j at time $t + 1$, given the model and the observation sequence.


$$\gamma_{t,i} = \gamma_t(i) = P(q_t = s_i | O, \lambda)$$

The probability of being in state s_i at time t , given the model and the observation sequence.

From the definitions it follows that:

$$\gamma_t(i) = \sum_{j=1}^N \xi_t(i,j)$$

Baum-Welch algorithm (II)



Outline

2 Machine Learning Applications for HMM

- Machine Learning
- Where do HMMs fit into Machine Learning?

3 Theory of HMMs

- The 3 things you want from an HMM
- Mathematical Foundations for HMMs
- Notation Conventions & Framework Description

4 Implementing HMMs

- Using the Model for Estimations: the Forward-Backward algorithm
- Learning from Observations: Baum-Welch algorithm
- Uncovering Hidden states: Viterbi algorithm

Viterbi s-a nascut in ...

Demo & Discussions

Outline

- 5 A Case for HMMs in Symbol Recognition
- 6 Types of HMMs
- 7 Discussions and Recap

Summary

- HMM are useful for temporal sequence problems

Summary

- HMM are useful for temporal sequence problems
- there are 3 fundamental problems for HMMs:

Summary

- HMM are useful for temporal sequence problems
- there are 3 fundamental problems for HMMs:
 - given an HMM, estimating the probability of an observed sequence

Summary

- HMM are useful for temporal sequence problems
- there are 3 fundamental problems for HMMs:
 - given an HMM, estimating the probability of an observed sequence
 - given observed data, estimating the parameters of an HMM

Summary

- HMM are useful for temporal sequence problems
- there are 3 fundamental problems for HMMs:
 - given an HMM, estimating the probability of an observed sequence
 - given observed data, estimating the parameters of an HMM
 - uncovering the hidden states

Summary

- HMM are useful for temporal sequence problems
- there are 3 fundamental problems for HMMs:
 - given an HMM, estimating the probability of an observed sequence
[Forward-Backward Algorithm](#)
 - given observed data, estimating the parameters of an HMM
 - uncovering the hidden states

Summary

- HMM are useful for temporal sequence problems
- there are 3 fundamental problems for HMMs:
 - given an HMM, estimating the probability of an observed sequence
[Forward-Backward Algorithm](#)
 - given observed data, estimating the parameters of an HMM
[Baum-Welch Algorithm](#)
 - uncovering the hidden states

Summary

- HMM are useful for temporal sequence problems
- there are 3 fundamental problems for HMMs:
 - given an HMM, estimating the probability of an observed sequence
[Forward-Backward Algorithm](#)
 - given observed data, estimating the parameters of an HMM
[Baum-Welch Algorithm](#)
 - uncovering the hidden states [Viterbi Algorithm](#)

References I

Dempster, A.P., N.M. Laird, and D.B. Rubin (1977). “Maximum likelihood from incomplete data via the EM algorithm”. In: *Journal of the Royal Statistical Society. Series B (Methodological)*, pp. 1–38.

Koller, D. and N. Friedman (2009). *Probabilistic Graphical Models: Principles and Techniques*. MIT Press.

Rabiner, L.R. (1989). “A tutorial on hidden Markov models and selected applications in speech recognition”. In: *Proceedings of the IEEE 77.2*, pp. 257–286.

Thank you!