Hidden Markov Models

Theoretical Aspects. Octave Implementation. Applications

Alexandru Sorici, Tudor Berariu

Romanian Asociation for Artificial Intelligence

October, 27th, 2012

PART 1.
Intro

PART 2.
Theory of HMMs

PART 3.
Demo & Discussions

Outline

Intro

Outline

Outline

1. ARIA Education Workshops
   - ARIA's Mission
   - ARIA Education
   - Workshop Program

Outline

ARIA EDU

:)

Outline

Today's Program

9:00	Registration
10:00	Rahaturi despre ARIA
11:00	HMM Theory

Theory of HMMs

Outline

2. Machine Learning Applications for HMM
   - Machine Learning
   - Where do HMMs fit into Machine Learning?

3. Theory of HMMs
   - The 3 things you want from an HMM
   - Mathematical Foundations for HMMs
   - Notation Conventions & Framework Description

4. Implementing HMMs
   - Using the Model for Estimations: the Forward-Backward algorithm
   - Learning from Observations: Baum-Welch algorithm
   - Uncovering Hidden states: Viterbi algorithm

Outline

What is Machine Learning?

Machine Learning

..Trascau is a beautiful horse.

Machine Learning Applications

- Computer Vision: Google Car
- Machine Translation
- Speech Recognition
- Recommender Systems
- Intelligent Advertising

Outline

Machine Learning Classification

Types of Machine Learning Problems

- Regression
- Classification
- Reinforcement Learning

- supervised learning (eg. ..)
- unsupervised

Sequence / Temporal problems (I)

OBJECT TRACKING

- Speaker Detection
- GPS
- Robotics
- Surface to air missille
- Ship or rocket navigation

SPEECH RECOGNITION

- Voice User Interfaces - e.g. SIRI
- Speech-to-Text Processing
- Direct Voice Input - Aircraft

GESTURE RECOGNITION

- Personalized Signature Recognition
- Human Activity Recognition
- Sign Language Recognition

Sequence / Temporal problems (II)

BIOINFORMATICS

Protein Sequencing

Modeling of a Gene
Regulatory Network

ECONOMICS

Stock Price Prediction

Econometrics
- estimate a country's econmic indicators across time -

Probabilistic Reasoning over Time - Models

Consider some of the previously presented problems ...

Probabilistic Reasoning over Time - Models

Consider some of the previously presented problems ...

How do we model such dynamic situations?

Probabilistic Reasoning over Time - Models

Consider some of the previously presented problems ...

How do we model such dynamic situations?

States and Observations

- The process of change is viewed as a series of time slices (snapshots)
- Each time slice contains a set of random variables
  - $\mathbf{O}_t$ - set of all *observable* evidence variables at time t
  - $\mathbf{Q}_t$ - set of all *unobservable / hidden* state variables at time t

Probabilistic Reasoning over Time - Assumptions

Consider some of the previously presented problems ...

Probabilistic Reasoning over Time - Assumptions

Consider some of the previously presented problems ...

What assumptions (if any) do we make?

Probabilistic Reasoning over Time - Assumptions

Consider some of the previously presented problems ...

What assumptions (if any) do we make?

Stationary Process

The process of change is governed by laws that do not themselves change over time.
Implication: we need to specify conditional distributions only for the variables within a *representative* timeslice.

Probabilistic Reasoning over Time - Assumptions

Consider some of the previously presented problems ...

What assumptions (if any) do we make?

Stationary Process

The process of change is governed by laws that do not themselves change over time.
Implication: we need to specify conditional distributions only for the variables within a *representative* timeslice.

Markov Assumption

The current state in a process of change depends only on a finite history of previous states.
Implication: there is a bounded number of "parents" for the variables in each time slice.

$\mathbf{P}(\mathbf{Q}_t|\mathbf{Q}_{1:t-1}) = \mathbf{P}(\mathbf{Q}_t|\mathbf{Q}_{t-1}) \qquad \mathbf{P}(\mathbf{O}_t|\mathbf{Q}_{1:t}, \mathbf{Q}_{1:t-1}) = \mathbf{P}(\mathbf{O}_t|\mathbf{Q}_t)$

Probabilistic Reasoning over Time - Inference

What are the basic inference tasks that must be solved?

Probabilistic Reasoning over Time - Inference

What are the basic inference tasks that must be solved?

Filtering (monitoring)

The task of computing the belief state - the posterior distribution over the current state, given all evidence to date.
$\mathbf{P}(\mathbf{Q}_t|\mathbf{o}_{1:t})$

Probabilistic Reasoning over Time - Inference

What are the basic inference tasks that must be solved?

Filtering (monitoring)

The task of computing the belief state - the posterior distribution over the current state, given all evidence to date.
$\mathbf{P}(\mathbf{Q}_t|\mathbf{o}_{1:t})$

Evaluation (likelihood)

The task of computing the likelihood of the evidence up to present.
$\mathbf{P}(\mathbf{o}_{1:t})$

Probabilistic Reasoning over Time - Inference

Prediction

The task of computing the posterior distribution over the future state, given all evidence to date.

$\mathbf{P}(\mathbf{Q}_{t+k}|\mathbf{o}_{1:t})$, for some $k > 0$

Probabilistic Reasoning over Time - Inference

Prediction

The task of computing the posterior distribution over the future state, given all evidence to date.
$\mathbf{P}(\mathbf{Q}_{t+k}|\mathbf{o}_{1:t})$, for some $k > 0$

Smoothing (hindsight)

The task of computing the posterior distribution over a past state, given all evidence to the present.
$\mathbf{P}(\mathbf{Q}_{k}|\mathbf{o}_{1:t})$, for some $1 \leq k < t$
Provides a better estimate of the state than was available at the time.

Probabilistic Reasoning over Time - Inference

Most likely explanation

Given a *sequence of observations*, find the sequence of states that is most likely to have generated those observations. $argmax_{q_{1:t}}\ \mathbf{P}(\mathbf{q}_{t+k}|\mathbf{o}_{1:t})$, for some $k > 0$

Probabilistic Reasoning over Time - Inference

Most likely explanation

Given a *sequence of observations*, find the sequence of states that is most likely to have generated those observations. $argmax_{q_{1:t}}\ \mathbf{P}(\mathbf{q}_{t+k}|\mathbf{o}_{1:t})$, for some $k > 0$

Learning

Given a set of *observation sequences*, find a method to learn the transition (e.g. $P(q_{t+1} = s_j|q_t = s_i)$, $1 \leq i, j < N$) and sensor ($P(o_t|q_t)$) models from the observations.

Probabilistic Reasoning over Time - Known Methods

Dynamic Bayesian Networks (DBN)

A DBN is Bayesian network that represents a temporal probability model.

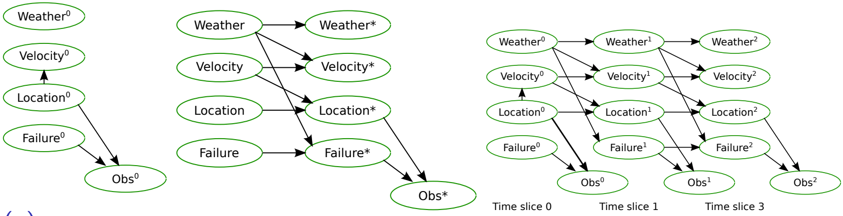

(a) the 2-time-slice Bayesian Network (b) the time 0 network (c) unrolled DBN over 3 time slices

Figure: A highly simplified DBN for monitoring a vehicle (Koller and Friedman 2009)

Applied in problems like: object tracking, human activity recognition, protein sequencing etc.

Probabilistic Reasoning over Time - Known Methods

Kalman Filters (Linear Dynamical Systems)

A temporal model of one or more real-valued variables that evolve linearly over time, with some Gaussian noise.

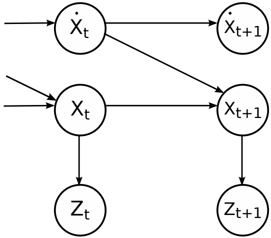

Figure: BN structure for a linear dynamical system with position X_t, velocity $\dot{X}_t$, and position measurement Z_t

- can be viewed as DBNs where all variables are continuous and all dependencies are linear gaussian
- wide application in **object tracking**

Probabilistic Reasoning over Time - Known Methods

Hidden Markov Models (HMM)

An HMM is a temporal probabilistic model in which the state of the process is described by a single discrete random variable. The possible values of the variable are the possible states of the world.

Used successfully in applications like:

- Handwriting Recognition
- Gesture Recognition
- Speech Recognition
- Part-of-Speech Tagging
- DNA Sequencing

Outline

The 3 fundamental problems (Rabiner 1989)

- Particularization of temporal inference problems to HMMs

The 3 fundamental problems (Rabiner 1989)

- Particularization of temporal inference problems to HMMs

Evaluation Problem

Given a model and a sequence of observations, how do we compute the probability that the observed sequence was produced by the model?

The 3 fundamental problems (Rabiner 1989)

- Particularization of temporal inference problems to HMMs

Evaluation Problem

Given a model and a sequence of observations, how do we compute the probability that the observed sequence was produced by the model?

Best Explanation of Observations Problem

Given a model and a sequence of observations how do we choose a corresponding sequence of states which *gives meaning* to the observations? How do we *uncover* the hidden part of the model?

The 3 fundamental problems (Rabiner 1989)

- Particularization of temporal inference problems to HMMs

Evaluation Problem

Given a model and a sequence of observations, how do we compute the probability that the observed sequence was produced by the model?

Best Explanation of Observations Problem

Given a model and a sequence of observations how do we choose a corresponding sequence of states which *gives meaning* to the observations? How do we *uncover* the hidden part of the model?

Model Estimation (Training) Problem

Given some observed sequences, how do we adjust the parameters of an HMM model that best tries to explain the observations?

Outline

Elements of an HMM

- N Hidden States : $S_1, S_2, \ldots S_N$
- M Observable Variables : $O_1, O_2, \ldots O_M$

Parameters:

- Transition Function / Matrix between states
- Emission probabilities
- Initial state probabilities

Formalisation of the estimation problem

-

Formalisation of problem # 2

- $P(Q_1) = \sum_{x=\{1,2\}}^{N} P(Q_2)\theta\Pi$
- $P(Q_i|q_i = s_x) = i \times x \cdot i \ldots$

Formalisation of parameters estimation problem

-

Outline

Notation Conventions

Variables in Octave

Outline

Outline

Outline

Viterbi s-a nascut in ...

Demo & Discussions

Outline

Summary

- HMM are useful for temporal sequence problems

Summary

- HMM are useful for temporal sequence problems
- there are 3 fundamental problems for HMMs:

Summary

- HMM are useful for temporal sequence problems
- there are 3 fundamental problems for HMMs:
  - given an HMM, estimating the probability of an observed sequence

Summary

- HMM are useful for temporal sequence problems
- there are 3 fundamental problems for HMMs:
  - given an HMM, estimating the probability of an observed sequence
  - given observed data, estimating the parameters of an HMM

Summary

- HMM are useful for temporal sequence problems
- there are 3 fundamental problems for HMMs:
  - given an HMM, estimating the probability of an observed sequence
  - given observed data, estimating the parameters of an HMM
  - uncovering the hidden states

Summary

- HMM are useful for temporal sequence problems
- there are 3 fundamental problems for HMMs:
  - given an HMM, estimating the probability of an observed sequence
    Forward-Backward Algorithm
  - given observed data, estimating the parameters of an HMM
  - uncovering the hidden states

Summary

- HMM are useful for temporal sequence problems
- there are 3 fundamental problems for HMMs:
  - given an HMM, estimating the probability of an observed sequence
    Forward-Backward Algorithm
  - given observed data, estimating the parameters of an HMM
    Baum-Welch Algorithm
  - uncovering the hidden states

Summary

- HMM are useful for temporal sequence problems
- there are 3 fundamental problems for HMMs:
  - given an HMM, estimating the probability of an observed sequence Forward-Backward Algorithm
  - given observed data, estimating the parameters of an HMM Baum-Welch Algorithm
  - uncovering the hidden states Viterbi Algorithm

References I

Koller, D. and N. Friedman (2009). *Probabilistic Graphical Models: Principles and Techniques*. MIT Press.

Rabiner, L.R. (1989). “A tutorial on hidden Markov models and selected applications in speech recognition”. In: *Proceedings of the IEEE* 77.2, pp. 257–286.

Thank you!